Warish Ali

Portfolio: warishali.com.np
Github: github.com/warishali1

Email: warishali.contact@gmail.com
Mobile: +977-9867779824

EDUCATION

Advanced College of Engineering and Management Kathmandu, Nepal
Bachelor in Computer Engineering Jan 2023 - Present
Courses: Artificial Intelligence, Probability & Statistics, Data Structures and Algorithms, Database Management Systems, Operating Systems, Computer Networks & Security

SKILLS SUMMARY

Languages & Databases: Python, SQL, Java, JavaScript

AI/ML Frameworks: NumPy, PyTorch, Scikit-Learn, Sentence Transformers

AI/ML: NLP, Transformers, Large Language Models (LLMs), RAG, Computer Vision, OCR

Full-Stack & Tools: Vue.js, Next.js, Spring Boot, MySQL, Git/GitHub, Postman

EXPERIENCE

Full Stack Java Developer Jan 2023 - Present
Raise Tech Pvt. Ltd.

- **Performance Improvements**: Optimized backend logic and database queries for core ERP modules including sales, purchase, invoice, and accounts, reducing page load and response times from ~ 5 minutes to under 1 minute while fixing major production bugs.
- **Security & Access Control**: Implemented core security and control features including OTP based login and password reset, admin-defined time-based access to the ERP and a notification system for payments and CRM followups.
- **Student CRM System**: Developed full-stack Student CRM system to track 100+ student's visa application progress end-to-end, including student management, progress tracking, and OTP authentication, on an existing production codebase.

PROJECTS

Applicant Tracking System June 2026

- Developed an ATS scoring engine using Sentence Transformers embeddings, scikit-learn similarity metrics, and rule-based ranking to score candidate-job fit across 4 weighted components, completing evaluations in an average of 18.8ms (44.2ms at the 95th percentile).
- Implemented a 3-tier domain-mismatch penalty that discounts scores by 7-20% when a candidate's skill profile diverges from a job's primary technical domain, paired with whole-word skill matching to eliminate false-positive substring matches.
- Built a resume ingestion pipeline that extracts text from PDF, DOCX, and plain text files with automatic OCR fallback for scanned documents, then uses an LLM to extract and normalize candidate information into a validated 15-attribute structured schema.

Exam Guide System April 2026

- Engineered an AI pipeline that processes uploaded notes, syllabus and past papers to generate a chapter-wise guide ranking each chapter by exam importance, with key topics, FAQs with frequency counts and marks distribution, reducing student study effort by up to 60%.
- Built a RAG pipeline that retrieves the most relevant passages from user-uploaded documents using semantic search and keyphrase extraction, feeding them as grounded context to an LLM to generate cited answers strictly from user content with no hallucinations.
- Developed a quiz engine that guarantees topic coverage by sampling questions across all chapters, ensuring every session surfaces new questions and unexplored topics through deduplication and rotation, generating 100+ unique questions per upload.

Transformer Language Model March 2026

- Built and trained a ~ 158M-parameter Transformer from scratch in PyTorch as an experimental project to explore attention mechanisms, encoder-decoder architecture, tokenization, decoding strategies, and training limitations under constrained compute (Google Colab).

Real-Time Face Recognition System June 2025

- Developed a real-time face recognition application using Python, Streamlit and PyTorch that detects and recognizes multiple faces within a second, allowing users to register new identities through uploaded images for fast and accurate recognition.
- Integrated pre-trained InceptionResnetV1 and MTCNN models for efficient face detection, facial alignment and feature extraction, enabling high-accuracy recognition through transfer learning without requiring computationally expensive model training from scratch.
- Built a user enrollment pipeline that stores face images and extracted embeddings in Google Drive for persistent storage across sessions, demonstrating practical implementation of CNNs, transfer learning and embedding-based recognition systems.